

Basic version of WineQualityMLP (Saved):

No normalization

Layers: [11, 5, 1]

Without dropout

Activation function: reLu


Loss function: MSE

Number of epochs:100

Batch size: 8

Learning rate: 0.0001

Loss curve:



The MSE of test set of the basic version of MLP: 0.5035768151283264

Tuning the model:

Hyperparameters can be tuned:

1. Normalization
 - a. Basic version: no normalization
 - b. Try normalize x (Standard Scaler works the best)
 - c. Try normalize both x values and y values
 - d. Try batch norm
2. Shuffle data when load
 - a. Basic version: no shuffle
 - b. Try with shuffle
1. Number of hidden layers and neurons per hidden layer
 - a. Basic version has [5]
 - b. Try [5, 5]
 - c. Try [10]
 - d. Try [10, 2]
2. Activation functions of each layer:
 - a. Basic version has relu
 - b. Try Sigmoid
 - c. Try Tanh
3. Number of epochs
 - a. Basic version has 100 epochs
 - b. Try early stopping
4. Learning rate:
 - a. 0.0001
 - b. 0.001
 - c. 0.01
5. Dropout rate:
 - a. 0
 - b. 0.1
 - c. 0.5
- 6.

Tune V1.1(Saved) :

Normalization on x: `StandardScaler()`

Layers: [11, 5, 1]

Without dropout

Activation function: reLu

Loss function: MSE


Number of epochs: 500

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss curve:



The MSE of test set: 0.42102354764938354

Tune V1.2:

Normalization on x: MinMaxScaler()

Layers: [11, 5, 1]

Without dropout

Activation function: reLu

Loss function: MSE


Number of epochs: 500

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss curve:



The MSE of test set: 0.4620991349220276

MinMaxScaler can slightly improve the performance, but not as good as the Standard Scaler.

Tune V1.3(Saved) :

Normalization on x: StandardScaler()

Layers: [11, 5, 5, 1]

Without dropout

Activation function: reLu

Loss function: MSE


Number of epochs: 500

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss curve:



The MSE of test set: 0.41365745663642883

TuneV1.4


Normalization on x: MinMaxScaler()

Layers: [11, 5, 5, 1]

Without dropout

Activation function: reLu
Loss function: MSE
Number of epochs: 500
Batch size: 8
Learning rate: 0.0001
Shuffle training data in dataloader

Loss curve:



The MSE of test set: 0.4536188840866089

This is similar to V1.2


TuneV1.5 (Saved)

Normalization: SS
Layers: [11, 10, 1]
Without dropout
Activation function: reLu
Loss function: MSE
Number of epochs: 500
Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss Curve:



The MSE of test set: 0.4028308093547821

TuneV1.6(Saved)

Normalization: SS

Layers: [11, 10, 10, 1]

Without dropout

Activation function: reLu

Loss function: MSE


Number of epochs:500

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss Curve:



The MSE of test set: 0.4162111282348633

TuneV1.7/V1.8

Normalization: SS

Layers: [11, 10, 10, 1]

Dropout_rate: 0.1/0.5

Activation function: reLu

Loss function: MSE


Number of epochs:500

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss Curve (V1.8):



(V1.7)The MSE of test set: 0.6362356543540955

(V1.8)The MSE of test set: 0.7410240173339844

Dropout will make the performance worse.

TuneV1.9/V1.10 (Saved)

Normalization: SS

Layers: [11, 10, 10, 1]

Dropout_rate: 0

Activation function: Sigmoid/Tanh

Loss function: MSE


Number of epochs:500

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss Curve(V1.10 Tanh):




(Sigmoid)The MSE of test set: 0.5983479619026184
(Tanh)The MSE of test set: 0.41961660981178284

TuneV2 (Selected from the best models from V1.*)

Normalization: SS
Layers: [11, 5, 5, 1]
Dropout_rate: 0
Activation function: Relu
Loss function: MSE
Number of epochs: 5000
Batch size: 8
Learning rate: 0.0001
Shuffle training data in dataloader

Loss curve:



The MSE of test set: 0.40147849917411804

All right, time to set up early stopping

TuneV2.1 (Saved)

Normalization: SS

Layers: [11, 5, 5, 1]

Dropout_rate: 0

Activation function: Relu

Loss function: MSE


Number of epochs: 5000 - early stop 100 (early stopped at epoch 2942)

Batch size: 8

Learning rate: 0.0001

Shuffle training data in dataloader

Loss Curve:



The MSE of test set: 0.3814166486263275

TuneV2.2/TuneV2.3

Normalization: SS

Layers: [11, 5, 5, 1]

Dropout_rate: 0

Activation function: Relu

Loss function: MSE

Number of epochs: 5000 - early stop 100

Batch size: 32/16

Learning rate: 0.0001

Shuffle training data in dataloader

(V2.2)The MSE of test set: 0.40


(V2.3)The MSE of test set: 0.40603408217430115

Increasing Batch size is not improving the proformance

TuneV2.4 (Saved)

Normalization: SS
Layers: [11, 5, 5, 1]
Dropout_rate: 0
Activation function: Relu
Loss function: MSE
Number of epochs: 5000 - early stop 100
Batch size: 8
Learning rate: 0.001
Shuffle training data in dataloader

Loss Curve:



The MSE of test set: 0.39196497201919556

TuneV2.5/TuneV2.6

Normalization: SS
Layers: [11, 5, 5, 1]
Dropout_rate: 0
Activation function: Relu

Loss function: MSE

Number of epochs: 5000 - early stop 100/20

Batch size: 8

Learning rate: 0.001

Shuffle training data in dataloader

momentum=0.5

(V2.5)The MSE of test set: 0.396

(V2.6)The MSE of test set: 0.38899001479148865

TuneV2.7(Saved)

Normalization: SS

Layers: [11, 5, 5, 1]

Dropout_rate: 0

Activation function: Relu

Loss function: MSE

Number of epochs: 5000 - early stop 20


Batch size: 8

Learning rate: 0.001

Shuffle training data in dataloader

momentum=0.2

Loss Curve




The MSE of test set: 0.37462902069091797

TuneV2.8

Normalization: SS
Layers: [11, 5, 5, 1]
Dropout_rate: 0
Activation function: Relu
Loss function: MSE
Number of epochs: 5000 - early stop 20
Batch size: 8
Learning rate: 0.001
Shuffle training data in dataloader
momentum=0.8

Loss curve:




The MSE of test set: 0.4131854176521301

TuneV2.9 (Saved)

Normalization: SS
Layers: [11, 5, 5, 1]
Dropout_rate: 0
Activation function: Relu
Loss function: MSE
Number of epochs: 5000 - early stop 50 (Stopping at epoch 929)
Batch size: 8
Learning rate: 0.0001
Shuffle training data in dataloader
momentum=0.2

Loss curve:



The MSE of test set: 0.38019827008247375

All right, I think the tuning is good enough!